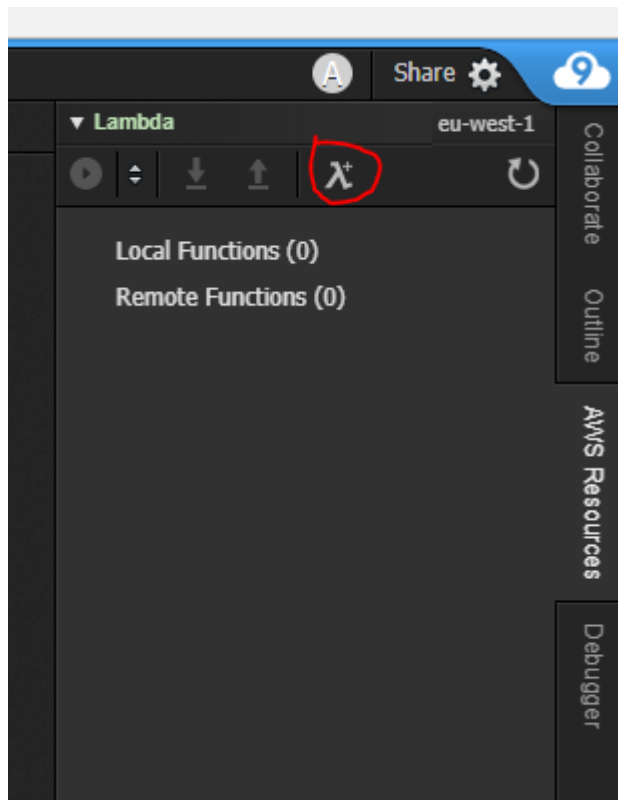




Building HelloWorld

Let's start off with organizing our work and creating a new folder for the mandatory "Hello World" example:

Create a new lambda function by using the AWS Resources menu and selecting the lambda+ icon:





Create serverless application

A serverless application can have one or more AWS Lambda functions, along with triggers and integrations for each of those functions. All of its configuration is stored in an AWS CloudFormation template file

Application name is used for the folder name for both your application's files and the AWS CloudFormation stack name when you deploy this application.

Function name:

Function name may only contain alphanumeric characters

Application name:

Application name may only contain alphanumeric characters

Region: eu-west-1

Next

The function needs to be developed. Let's selected Python 3.6+ as the run time to write the Hello World example as Python code:



Create serverless application

Select runtime

All runtimes

Select blueprint

empty-nodejs An empty NodeJS function nodejs · nodejs6.10	empty-python An empty Python function python · python3.6
alexa-skill-kit-sdk-factskill Demonstrate a basic fact skill built with the ASK NodeJS SDK nodejs6.10 · alexa	alexa-skills-kit-color-expert Demonstrates a basic skill built with the Amazon Alexa Skills Kit. nodejs6.10 · alexa
alexa-skills-kit-color-expert Demonstrates a basic skill built with the Amazon	alexa-smart-home-skill-adapter Provides the basic framework for a skill adapter for a

Region: eu-west-1

Previous Next

For now, we don't require a function trigger (we will test the function manually first):

Create serverless application

Function trigger

none

Region: eu-west-1

Previous Next

Also, default memory settings and a default generated role should be enough - we're not doing any heavy lifting, nor are we using any asynchronous functionality:



Create serverless application

Memory (MB): 128 MB

Your function is allocated CPU proportional to the memory configured.

Role: Automatically generate role

Region: eu-west-1

Previous Next

Review the settings and finish:

Create serverless application

Review your serverless application details below. You can go back to make changes for each section. When you are ready, click Finish to complete the setup process.

After AWS Cloud9 creates your new application, it will be automatically deployed using AWS CloudFormation.

Lambda function

Region	eu-west-1
Name	HelloWorld
ApplicationName	HelloWorld
Runtime	python3.6
Blueprint	empty-python
Handler	lambda_function.lambda_handler
Memory-size	128
Role	Automatically Generated

Region: eu-west-1

Previous Finish

From the setting above, you should note that the Handler is “`lambda_function.lambda_handler`”. What does this mean? Well, if the function is invoked, it will look for a Python **source file** named “`lambda_function`” (.py, actually) and try and execute the **function** “`lambda_handler`” inside this file. To isolate this project from other projects in the same IDE, Cloud9 will create a so-called “Virtual Environment” for Python, where all its dependencies will be stored and not pollute the global namespace - reducing the chances of conflicts.

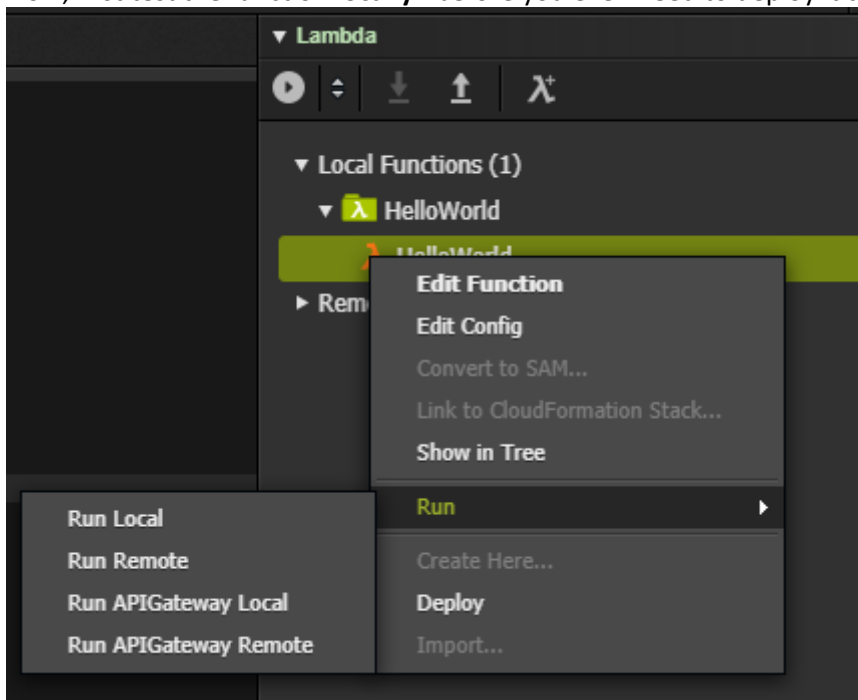
As you can see, the AWS Lambda framework will provide the function with an event and a context: the event is the event that triggered the invocation of the function, the context consists of AWS



specific environment information, the so-called execution context. Let's examine both objects passed in and return a greeting using "Hello " plus the name if present in the input event, otherwise name should be defaulting to "World" ".

```
1 import json
2
3 print('Loading...')
4
5 def lambda_handler(event, context):
6     # event is a dict object
7     print('Got an event: ' + json.dumps(event, indent=2))
8
9     # context is an object - https://docs.aws.amazon.com/lambda/latest/dg/python-context-object.html
10    print("Execution context - {} ms remaining".format(context.get_remaining_time_in_millis()))
11    print('Execution context: function {} version {}, size {} mb'.format(context.function_name, context.function_version, context.memory_limit_in_mb))
12    print('Logging to log group {} and log stream {}'.format(context.log_group_name, context.log_stream_name))
13
14    naam = 'World' if 'name' not in event else event['name']
15
16    groet = 'Hello ' + naam + "!"
17    print('We gaan groeten ... ' + groet)
18
19    return groet
20
21 print('Done loading!')
```

Now, first test the function **locally** - before you even need to deploy it to the real cloud:



This opens a window that lets you specify an **event payload**. As we are expecting a name property at a minimum set this property in the event payload:



```

bash - "ip-172-31-16" x [A] HelloWorld x +
Run /HelloWorld/.debug/HelloWorld/lambda_function.py

▼ Test payload
Function: HelloWorld
Payload:
1 {
2   "name" : "SynTouch Colleagues",
3   "array" : [ "This", "is", "a", "structured", "object"],
4   "error" : false
5 }

▼ Execution results
Response
"Hello SynTouch Colleagues!"

Function Logs
Loading...
Done loading!
Got an event: {
  "name": "SynTouch Colleagues",
  "array": [
    "This",
    "is",
    "a",
    "structured",
    "object"
  ],
  "error": false
}
Execution context - 15000 ms remaining
Execution context: function test version $LATEST, size 128 mb
Logging to log group /aws/lambda/test and log stream 2018/7/4/[$LATEST]ffda7bf63c5909fe

```

When supplying a name attribute in the JSON event, the contents are used for the greeting. You can see the function loading (initializing), finish the initialization and then executing on the supplied event. As we're running locally, the actual names ("test" for the function name) may be different from the real run-time environment values.

What Cloud9 is doing behind the scenes, is creating a SAM (Serverless Application Model) template to provide all resources required. Actually, this is an extension to the CloudFormation service, which creates or updates an application stack based on an input template ('infrastructure' as code!)

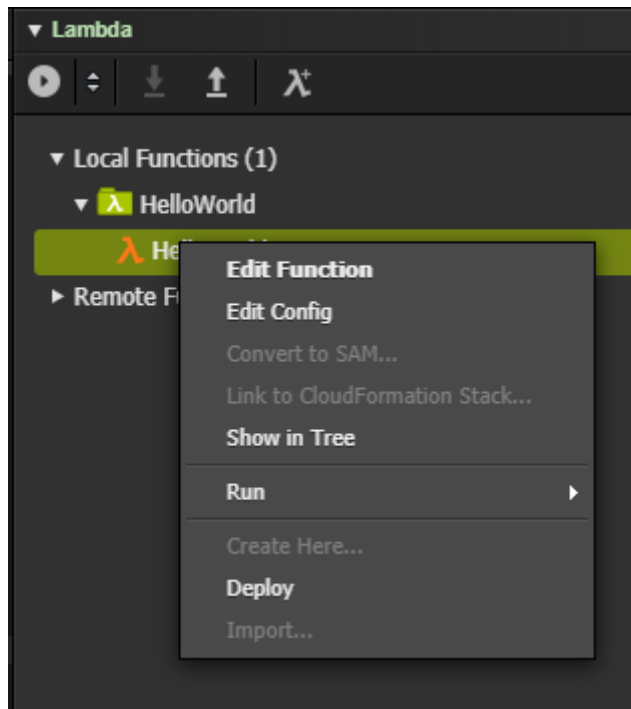
Before deployment, we need to make a small change to the template that is used to provision the stack (if we don't we cannot control the actual name assigned to the function completely): add a `FunctionName` attribute to the HelloWorld function and set this to HelloWorld:

```

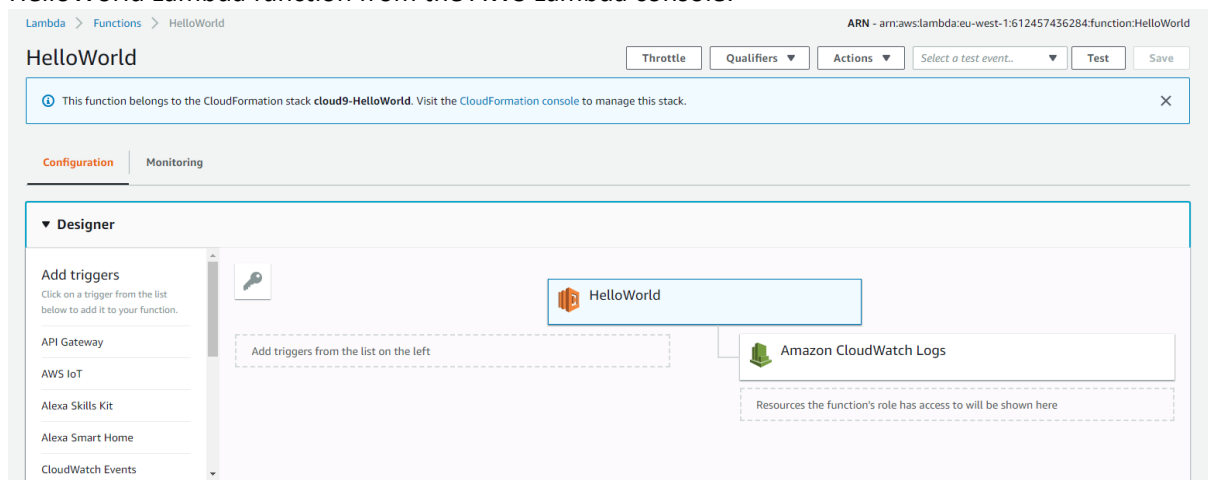
lambda_function.py x template.yaml x +
1 AWSTemplateFormatVersion: '2010-09-09'
2 Transform: 'AWS::Serverless-2016-10-31'
3 Description: An AWS Serverless Specification template describing your function.
4 Resources:
5   HelloWorld:
6     Type: 'AWS::Serverless::Function'
7     Properties:
8       Handler: HelloWorld/lambda_function.lambda_handler
9       FunctionName: HelloWorld
10      Runtime: python3.6
11      Description: 'This is my super duper HelloWorld function!'
12      MemorySize: 128
13      Timeout: 15
14      CodeUri: .debug/
15

```

Now the code is ready to rock & roll!



After some seconds, a stack has been provisioned using CloudFormation. You can examine the actual HelloWorld Lambda function from the AWS Lambda console:



Examine and edit the code:



HelloWorld Throttle Qualifiers ▼ Actions ▼ Select a test event.. ▼ Test Save

Function code [Info](#)

Code entry type: Edit code inline ▼ Runtime: Python 3.6 ▼ Handler: HelloWorld/lambda_function.lambda_handler [Info](#)

File Edit Find View Goto Tools Window

Environment

- ▼ HelloWorld
 - HelloWorld
 - ip3db-1.1.4.dist-info
 - __init__.py
 - easy_install.py
 - ip3db.py
 - isette3.py
 - requirements.txt

lambda_function

```

1 import json
2 print('Loading...')
3
4 def lambda_handler(event, context):
5     # event is a dict object
6     print('Got an event: ' + json.dumps(event, indent=2))
7
8     # context is an object - https://docs.aws.amazon.com/lambda/latest/dg/python-context-object.html
9     print('Execution context: ' + json.dumps(context.get_remaining_time_in_millis()))
10    print('Execution context: function {} version {}, size {} mb'.format(context.function_name, context.function_version, context.memory_limit_in_mb))
11    print('Logging to log group {} and log stream {}'.format(context.log_group_name, context.log_stream_name))
12
13    name = 'World' if 'name' not in event else event['name']
14
15    greet = 'Hello ' + name + '!'
16    print('We greet you ... ' + greet)
17    return greet
18
19 print('Done loading!')
```

1:1 Python Spaces: 4

And inspect other relevant settings:

Execution role

Defines the permissions of your function. Note that new roles may not be available for a few minutes after creation. [Learn more](#) about Lambda execution roles.

Choose an existing role ▼

Existing role

You may use an existing role with this function. Note that the role must be assumable by Lambda and must have CloudWatch Logs permissions.

cloud9-HelloWorld-HelloWorldRole-11HH2LFWGK063 ▼

Basic settings

Description

This is my super duper HelloWorld function!

Memory (MB) [Info](#)

Your function is allocated CPU proportional to the memory configured.

128 MB

Timeout [Info](#)

0 min 15 sec

The role generated by IAM allows the function to only write logging to CloudWatch:

Summary Delete role

Role ARN: [arn:aws:iam::612457436284:role:cloud9-HelloWorld-HelloWorldRole-11HH2LFWGK063](#)

Role description: [Edit](#)

Instance Profile ARNs: [Edit](#)

Path: /

Creation time: 2018-07-04 20:59 UTC+0200

Maximum CLI/API session duration: 1 hour [Edit](#)

Permissions Trust relationships Access Advisor Revoke sessions

Attach policies Attached policies: 1 [Add inline policy](#)

Policy name: [AWSLambdaBasicExecutionRole](#) Policy type: [AWS managed policy](#)

Policy summary { } JSON [Simulate policy](#)

```

1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Action": [
7         "logs:CreateLogGroup",
8         "logs:CreateLogStream",
9         "logs:PutLogEvents"
10      ],
11      "Resource": "*"
12    }
13  ]
14 }
```

And only AWS Lambda is allowed to assume this role:



Roles > cloud9-HelloWorld-HelloWorldRole-11HH2LFWGK063

Summary

Role ARN	arn:aws:iam::612457436284:role/cloud9-HelloWorld-HelloWorldRole-11HH2LFWGK063
Role description	Edit
Instance Profile ARNs	
Path	/
Creation time	2018-07-04 20:59 UTC+0200
Maximum CLI/API session duration	1 hour Edit

Permissions Trust relationships Access Advisor Revoke sessions

You can view the trusted entities that can assume the role and the access conditions for the role. [Show policy document](#)

[Edit trust relationship](#)

Trusted entities

The following trusted entities can assume this role.

Trusted entities

The identity provider(s) `lambda.amazonaws.com`

Conditions

The following conditions define how and when trusted entities can assume the role.

There are no conditions associated with this role.

Return to the AWS HelloWorld Lambda function and create a new test event by using the Test button:

Configure test event

A function can have up to 10 test events. The events are persisted so you can switch to another computer or web browser and test your function with the same events.

☒ Create new test event

☐ Edit saved test events

Event template

Hello World

Event name

MyHelloWorldEvent

```

1 {
2   "name": "SynTouch Colleagues",
3   "array": ["This", "is", "a", "structured", "object"],
4   "error": false
5 }
6

```

Save the event and test!

Lambda > Functions > HelloWorld

ARN - arn:aws:lambda:eu-west-1:612457436284:function:HelloWorld

HelloWorld

Throttle Qualifiers Actions MyHelloWorldEvent Test Save

This function belongs to the CloudFormation stack `cloud9-HelloWorld`. Visit the [CloudFormation console](#) to manage this stack.

The execution logs provide you with the actual output, but also with useful information about the duration, the duration you will be billed for (rounded up to 100 ms) and the actual memory consumed:



Lambda > Functions > HelloWorld

ARN - arn:aws:lambda:eu-west-1:612457436284:function:HelloWorld

Throttle Qualifiers Actions MyHelloWorldEvent Test Save

Execution result: succeeded (logs)

▼ Details

The area below shows the result returned by your function execution.

```
"Hello SynTouch Colleagues!"
```

Summary

Code SHA-256	+bmQ55md2ISY6IFTWtR6pwyVx1WnQN7ZEKHKPQWtNB8=	Request ID	2f13072f-7fbc-11e8-82cf-df068de8b3bb
Duration	0.64 ms	Billed duration	100 ms
Resources configured	128 MB	Max memory used	22 MB

Log output

The area below shows the logging calls in your code. These correspond to a single row within the CloudWatch log group corresponding to this Lambda function. [Click here](#) to view the CloudWatch log group.

```

{
  "object":
},
  "error": false
}
Execution context - 14999 ms remaining
Execution context: function HelloWorld version $LATEST, size 128 mb
Logging to log group /aws/lambda/HelloWorld and log stream 2018/07/04/[$LATEST]d9c71e52eb4c4bdf849da2228bdc4
We gaan groeten ... Hello SynTouch Colleagues!
END RequestId: 2f13072f-7fbc-11e8-82cf-df068de8b3bb
REPORT RequestId: 2f13072f-7fbc-11e8-82cf-df068de8b3bb Duration: 0.64 ms Billed Duration: 100 ms Memory Size: 128 MB Max Memory Used: 22 MB
  
```

The monitoring tab returns miscellaneous statistics on your function, but also offers access to the cloudwatch log files where this time period was logged:

Configuration **Monitoring**

CloudWatch metrics at a glance

Invocation count Last 24 hours ▼

[Jump to Metrics](#) [Jump to Logs](#)

Count (Double-click on the graph to zoom out)

CloudWatch > Log Groups > aws/lambda/HelloWorld > All streams

Expand all Filter events

Time (UTC +00:00)	Message	Stream
2018-07-04	No data events found for the selected date range. Adjust the date range.	
2018-07-04	Logging	/aws/lambda/HelloWorld
2018-07-04	Code Upload	/aws/lambda/HelloWorld
2018-07-04	START: RequestId: 2f13072f-7fbc-11e8-82cf-df068de8b3bb version: \$LATEST	/aws/lambda/HelloWorld
2018-07-04	END RequestId: 2f13072f-7fbc-11e8-82cf-df068de8b3bb	/aws/lambda/HelloWorld
2018-07-04	REPORT RequestId: 2f13072f-7fbc-11e8-82cf-df068de8b3bb Duration: 0.64 ms Billed Duration: 100 ms Memory Size: 128 MB Max Memory Used: 22 MB	/aws/lambda/HelloWorld

Expose HelloWorld as a WebAPI

Now that we have defined the mandatory HelloWorld as a serverless function in AWS Lambda, let's go the extra mile and expose this to the GBI ("Grote Boze Internet") as an WebAPI to be called over HTTP.

Make sure that you're still in the eu-west-1 (Ireland) region, as this is where your Lambda function lives.

Define a new API in API Gateway:



Amazon API Gateway APIs > Create

Create new API

In Amazon API Gateway, an API refers to a collection of resources and methods that can be invoked through HTTPS endpoints.

☒ New API ☐ Import from Swagger ☐ Example API

Settings

Choose a friendly name and description for your API.

API name*

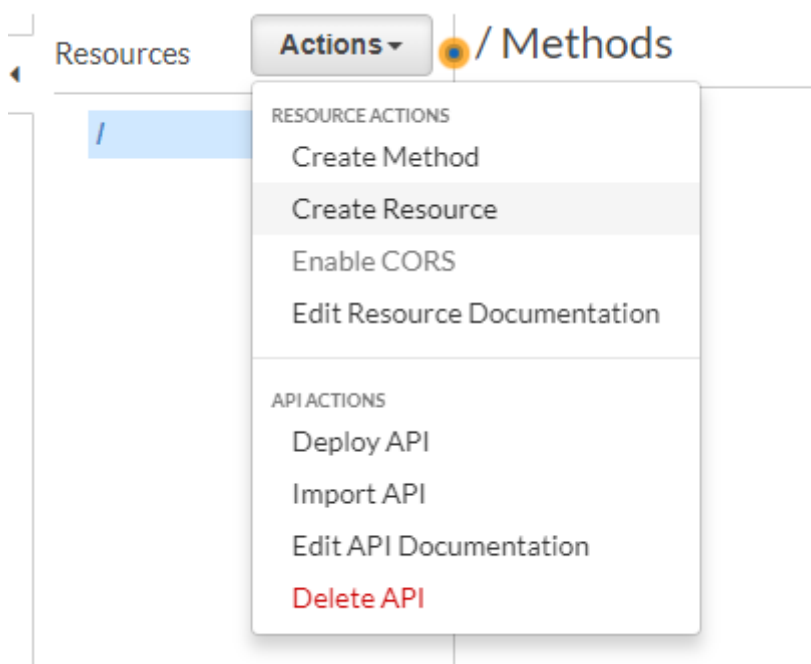
Description

Endpoint Type ⓘ

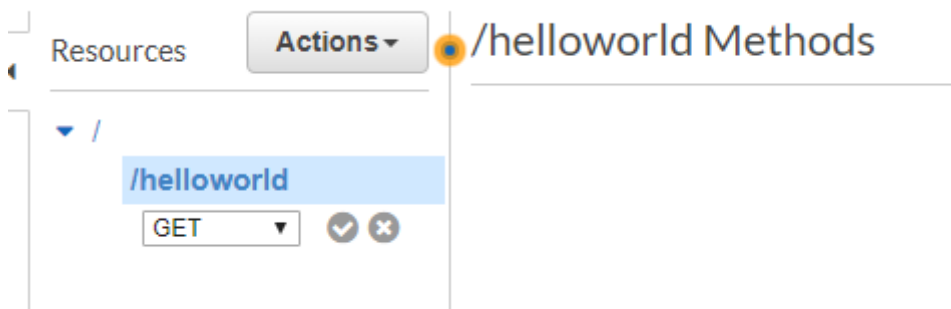
* Required

We'll define the API as a GET on resource helloworld where the name is to be provided as a query parameter ...

From the drop down, choose "Create Resource":



Select the /helloworld resource and choose "Create Method" from the dropdown list; select the GET HTTP method:





If you press the “Okay” icon, you’re presented with a configuration form to connect your exposed endpoint. Of course we’ll be connecting this to the HelloWorld lambda function:

Resources Actions /helloworld - GET - Setup

Choose the integration point for your new method.

Integration type ☒ Lambda Function ⓘ
☐ HTTP ⓘ
☐ Mock ⓘ
☐ AWS Service ⓘ
☐ VPC Link ⓘ

Use Lambda Proxy integration ☐ ⓘ

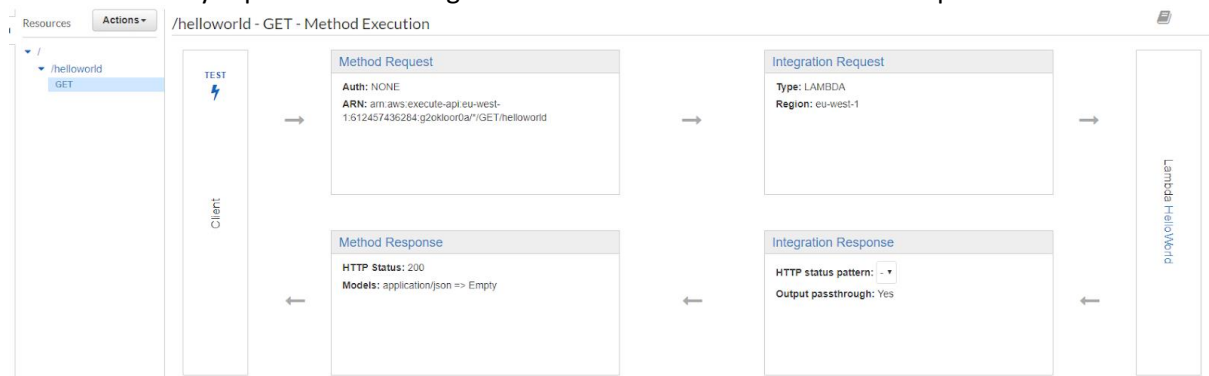
Lambda Region eu-west-1

Lambda Function HelloWorld ⓘ

Use Default Timeout ☒ ⓘ

When asked, authorize API Gateway to invoke your Lambda function.

The API Gateway represents the integration with the backend as a number of phases:



In the first stage (Method Request), the expected payload structure, headers and query string parameters can be defined. The request can be rejected if it does not satisfy the defined requirement. For our current purpose, this is not relevant.

The next step, Integration Request, is used to transform the received request to the request the backend (in this case our HelloWorld lambda function) will receive.

Here, we need to find a way to transfer the contents of the name querystring parameter to a name element in the JSON payload for the backend service. API Gateway supports the Apache Velocity templating engine for performing these tasks. For an overview of the available mapping functions, see <https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-mapping-template-reference.html>

Define a new template to be used for application/json content types; insert your expression replacing your-mapping-expression-here-to-transfer-the-querystring-param-name-value:



Resources **Actions** ▾

- ▾ /
 - ▾ /helloworld
 - GET

▸ URL Query String Parameters

▸ HTTP Headers

▾ Mapping Templates

Request body passthrough ☒ When no template matches the request Content-Type header ⓘ ⚠ ☐ When there are no templates defined (recommended) ⓘ ☐ Never ⓘ

Content-Type

application/json ⓘ

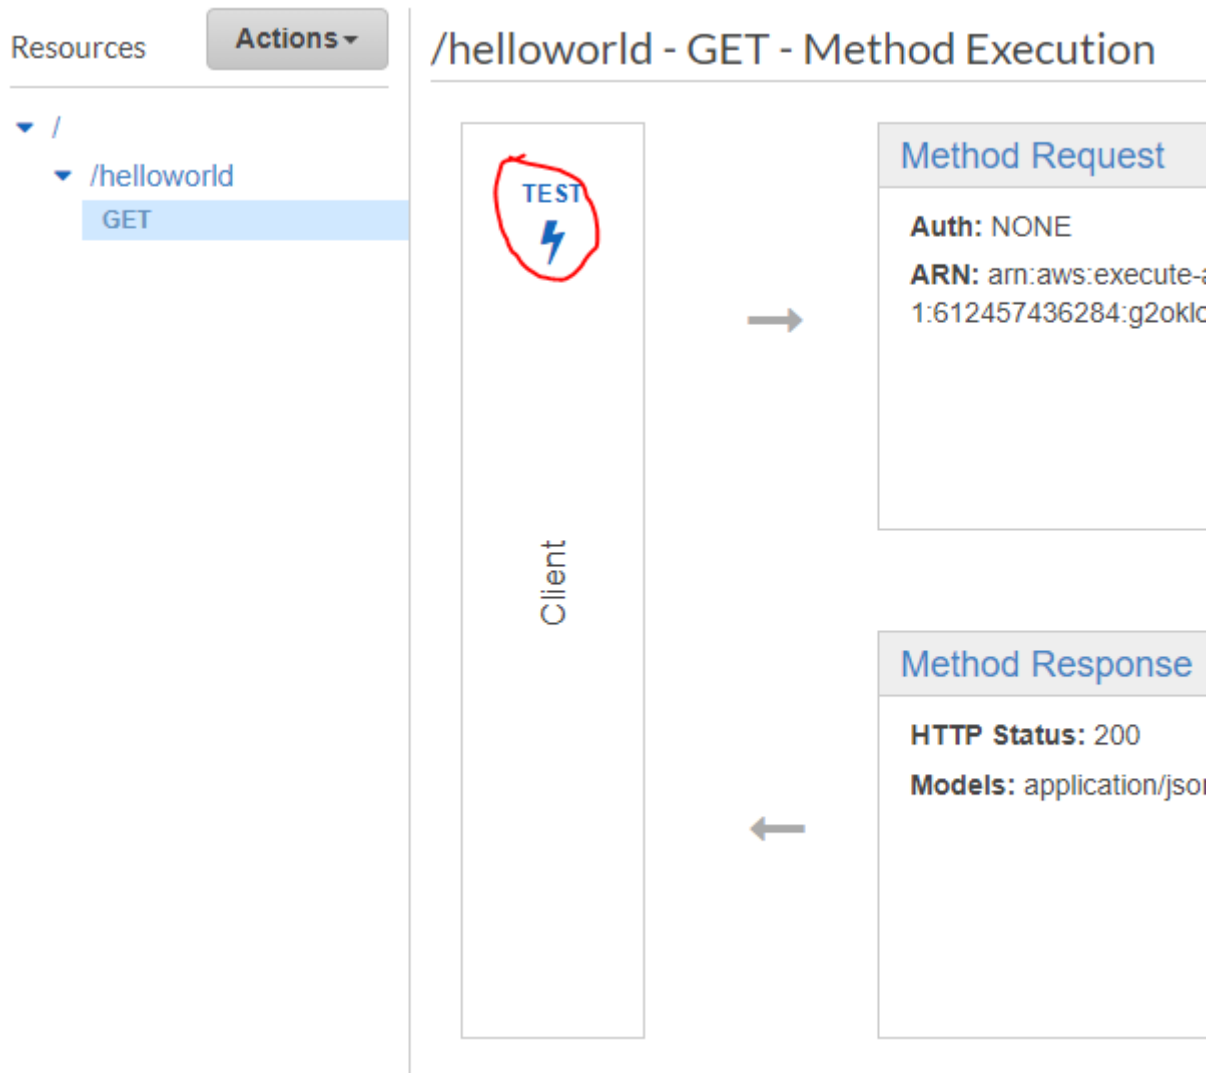
⊕ Add mapping template

application/json

Generate template: ▾

```
1 { "name" : your-mapping-expression-here-to-transfer-the-querystring-param-nam-value }
2
```

Navigate back to your API, select the GET method on helloworld and invoke the TEST button on the client:



Next, provide the query string parameter name and set to a value of your choice. TEST!

The screenshot shows the AWS Lambda console's 'Method Execution' page for the `/helloworld` GET method. The 'Query Strings' section has a text input with the value `name=World And Your Mama`. The 'Response Body' section shows the response `"Hello World And Your Mama!"`. The 'Logs' section shows the execution log for the request, including the request path, query string, and headers.

As you can see from the logs (and the response), the query string parameter is mapped to the `name` element in the payload of the event. However, the response is still returned as plain old text - since



the function simply returns text. This can be easily solved using the Integration Response mapping template, in a similar fashion we have applied to the request:

Resources Actions ▾

← Method Execution /helloworld - GET - Integration Response

First, declare response types using Method Response. Then, map the possible responses from the backend to this method's response types.

Lambda Error Regex	Method response status	Output model	Default mapping
-	200		Yes

Map the output from your Lambda function to the headers and output model of the 200 method response.

Lambda Error Regex

Content handling

Cancel Save

Header Mappings

Mapping Templates

Content-Type
application/json

Add mapping template

Add integration response

Add a mapping template to take the entire response body and assign this to the JSON message element; the resulting response should be like:

```
{ "message" : "The-actual-response-from-your-lambda-should-go-here..." }
```

Test the GET method again:



Resources

Actions ▾

▼ /

▼ /helloworld

GET

← Method Execution

/helloworld - GET - Method Test

Make a test call to your method with the provided input

Path

No path parameters exist for this resource. You can define path parameters by using the syntax **{myPathParam}** in a resource path.

Query Strings

{helloworld}

name=Lionel Richie

Headers

{helloworld}

Use a colon (:) to separate header name and value, and new lines to declare multiple headers. eg.
Accept:application/json.

Stage Variables

No [stage variables](#) exist for this method.

Request Body

Request Body is not supported for GET methods.

⚡ Test

Now is properly responding JSON !

```
{
  "message": "Hello Lionel Richie!"
}
```

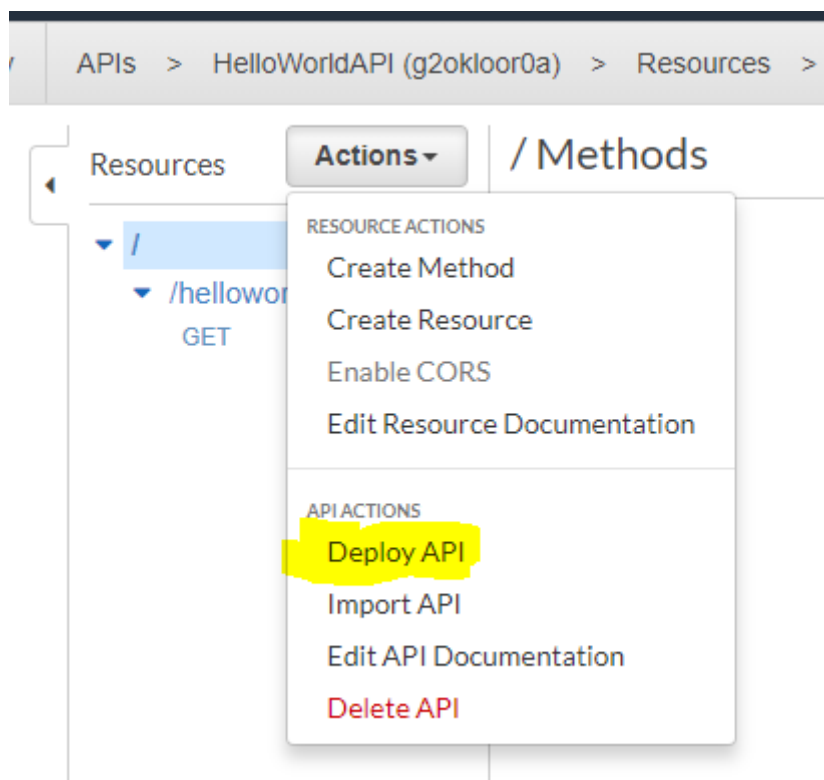
Response Headers

```
{"X-Amzn-Trace-Id": "Root=1-5b4a5ba8-7adcdff64f418131abe48e20;Sampled=0", "Content-Type": "application/json"}
```

Check the Logs for the steps executed upon receiving the request!

Making the API accessible

The final step in the process is actually deploying the API to the world:



Specify a stage (prefix) for your API and perform the deploy:

Deploy API

Choose a stage where your API will be deployed. For example, a test version of your API could be deployed to a stage named beta.

Deployment stage

[New Stage]

Stage name*

dev

Stage description

Development

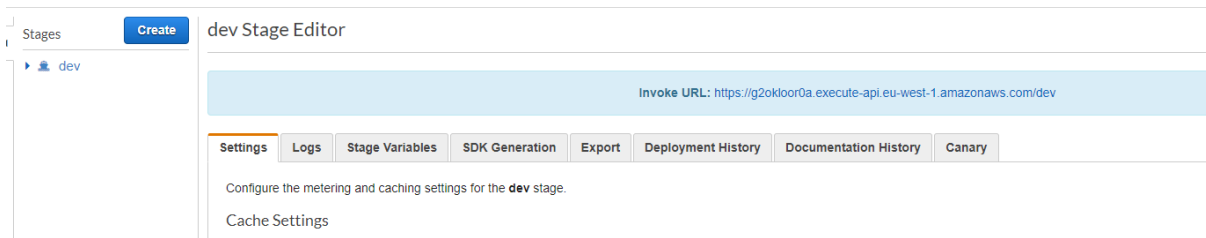
Deployment description

Vooruit met de geit!

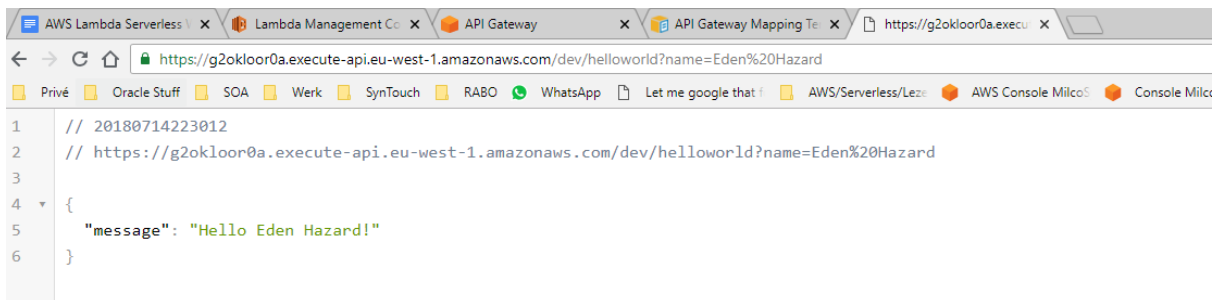
Cancel

Deploy

Examine the confirmation:



As you can see, here the stage “dev” is exposed at “<https://g2okloor0a.execute-api.eu-west-1.amazonaws.com/dev>”. The resource for our API is helloworld and the GET request requires a name parameter, so you can test the api by going to <https://g2okloor0a.execute-api.eu-west-1.amazonaws.com/dev/helloworld?name=Eden%20Hazard>



Well done!